

Threads

Introduction

Interruptions

Synchronization

Atomic

volatile

Wait/Notify

Locks and Monitors

Executors and Thread Pools

Immutable Objects

Thread-Safe Collections

INTRODUCTION

Threads

- A *thread* is a separate computation process that might execute in parallel with other processes.
- You use threads all the time:
 - Having a video playing while you are typing an email
 - Reading a website while you have music playing
- Java programs can have multiple threads.

Java Memory: Heap and Stack

- Stack Memory
 - Local variables and object references are stored in stack memory
 - Method calls (*activation records*) are stored on the stack
 - Variables on the stack exist as long as the method that created them is running
- Heap Memory
 - The actual objects are stored on the heap
 - The *reference* points to the location on the heap
- Good article/pic here: <https://www.journaldev.com/4098/java-heap-space-vs-stack-memory>

Java Threads

- Each thread has its own stack memory:
 - method activation records
 - local variables
 - parameters
- But threads share the heap.
 - objects are shared!
- Threads share resources (variables and data).
 - This happens so fast that it appears that the threads are operating in parallel.
- This is more efficient and allows for easy communication...
- But of course it's risky to have shared resources.

Single Threaded Programs

- All Java programs use threads.
- By default, each program uses a single thread.
 - The *main thread*
- If you've never used threads directly... then you've been writing one-thread programs all along!

static Thread Methods

- Thread.sleep
 - Pauses the execution of the current thread for a specified number of milliseconds (ish)
 - Throws a checked exception (InterruptedException) that must be caught (or propagated)
- Thread.activeCount()
- Thread.enumerate(Thread[] activeThreadArray)
- Thread.currentThread()
- Can set a name on a Thread using setName(name) or setName method, which is useful when distinguishing between threads

Practice

- Review the example that invokes methods on the main thread.

CREATING THREADS

Creating Your Own Threads

- You can create and manage threads in your program
- Two ways: extend Thread or implement Runnable

The Thread Class

- Your class extends `Thread`
 - Overrides the `public void run()` method
 - Can call `super(name)` to set a name for the thread

- You start by invoking the start method:

```
myThreadObject.start()
```

- **Do not call run- call start.**
- The start method has setup code it needs to execute before it invokes run.
- Invoking run doesn't actually create a new thread!!

The Runnable Interface

- Your class implements `Runnable`
 - Overrides the `public void run()` method
- You start by sending an instance of your class to the `Thread` constructor and invoking `start`:

```
Thread thread = new Thread(myRunnableObject)
thread.start()
```

- Again, **do not call run- call start**.
- The `start` method has setup code it needs to execute before it invokes `run`.
- Invoking `run` doesn't actually create a new thread!!

The Runnable Interface

- Runnable is a functional interface!
 - Queue the lambda entrance music!
- If you have a simple run method, you can skip creating an entire class that implements Runnable and just use a lambda:

```
new Thread(  
    () -> myRunningMethod()  
).start()
```

Practice

- Review the program that prints dates using Thread and Runnable.

When should I use concurrency?

- Tasks that will take a long time to execute.
- Any program where you want one thing to be able to run in the background and do other things too.
- GUI programs- you want the program to remain responsive while it is performing tasks.
 - **NOTE:** You cannot use the Thread and Runnable setup with JavaFX.
 - The scene graph is not thread-safe and can only be accessed and modified from the UI thread called the JavaFX Application thread.
 - The javafx.concurrent package supposed multithreaded GUIs.

INTERRUPTIONS

Ending Threads

- A thread terminates when the run method terminates by:
 - returning,
 - reaching the end, or
 - throwing an uncaught exception
- You can also request a thread to terminate with an *interrupt*
 - You can't force it! But you can request it.

Thread Interruptions

- Each thread has a boolean flag that represents whether it is interrupted

Requesting a Thread Interruption

- `void interrupt()`
 - Sends an interrupt request to the thread
 - The flag of the thread is set to true
 - If the thread is *blocked* (it is sleeping or waiting), an `InterruptedException` is thrown

Responding to Thread Interruptions

- A thread *could* just ignore interruptions! But this isn't convention.
- Threads should occasionally check to see if they've been interrupted.
 - This code is not automatic- you have to write it!
- It's common to stop a thread when it's been interrupted.
 - But again, this isn't automatic!

Responding to Thread Interruptions

- **static** boolean interrupted()
 - Tests whether the *current* thread has been interrupted
 - Side effect: resets the flag to false (clears the interruption status)
- boolean isInterrupted()
 - Tests whether a thread has been interrupted (**invoke on a Thread object**)
 - No side effects
 - Often use with: Thread.currentThread()
 - Thread.currentThread().isInterrupted()

Supporting Interruptions

- Inside a run that uses sleeps:

```
try {  
    while (moreWorkToDo) {  
        // do the work  
        // sleep  
    }  
} catch (InterruptedException ex) {  
    // thread interrupted during sleep/wait  
    return;  
}
```

Supporting Interruptions

- Inside a run that doesn't sleep:

```
while (!Thread.currentThread().isInterrupted())
    && moreWorkToDo)
    // do the work
}
// exit the run method
```

Supporting Interruptions

- Inside a run that doesn't sleep:

```
try {
    while (moreWorkToDo) {
        if (Thread.currentThread().isInterrupted()) {
            // throw new InterruptedException();
        } else {
            // do the work
        }
    }
} catch (InterruptedException ex) {
    // react to all interruption (perhaps return)
}
// exit the run method
```


Practice

- Interrupt the time print thread.
 - When it's asleep
 - While it's running
 - Change the behavior of how the thread responds to the interrupt

Join

- Invoking join makes the current thread wait until another thread finishes to continue.
- Example: invoking `myThread.join()` from main will make the main method wait until `myThread` finishes execution.

Practice

- Wait until the time programs are done to print a final message.

SYNCHRONIZATION

Shared Resources

- Multiple threads share access to the same data.
 - Objects are on the shared heap!
- If two threads each modify the state of the object, unpredictable or incorrect results can occur.
- This is called a *race condition*.

Race Conditions

- Race conditions can occur on any statement that is not *atomic* (meaning *happening all at once*).
- Even simple statements have multiple steps:
 - `i++`
 - `i = i + 1;`
 - get the current value of `i`
 - add one to it
 - store the new value of `i` back to the variable
- You could have different threads working at different steps!

Practice

- Review the simple counter example.
 - Counter, RaceConditionTest
- Review the bank account example (bonus/fee).
 - BankAccount, BankThread, BankTester

Synchronized

- The most straightforward way to prevent a race condition is to *synchronize*.
- Synchronizing allows only one thread to access a block of code at a time.
- You can synchronize an entire method:

```
public synchronized void myMethod() {  
    ... critical code  
}
```
- Or a block of code:

```
public void myMethod() {  
    ... non critical code  
    synchronized(someObject) { // could be "this"  
        ... critical code  
    }  
}
```
- You can also synchronize static methods.

Synchronized

- Note that this:

```
public synchronized void myMethod() {  
    ... critical code  
}
```

- Is equivalent to this:

```
public void myMethod() {  
    synchronized(this) {  
        ... critical code  
    }  
}
```

Synchronized

- When a thread reaches synchronized code, the JVM checks to see if anyone holds the lock **associated with the sync object**.
 - If no one does, the current thread gets the lock.
 - If someone else does, the current thread is placed in the waiting area.
- When a thread is done with synchronized code, it releases the lock.
 - The JVM then checks the waiting area to see if someone else needs the lock.
- Synchronizing works great... but it can lead to inefficiency, so it should be used cautiously.

Practice

- Update the bank example to avoid race conditions.

Shortcomings of synchronized

- It can be easy to make mistakes that make code appear to be synchronized, when it's really not
 - To avoid this, make sure to choose an object common to all threads.
- Deadlock
 - Avoid this by being careful about how synchronized methods invoke other synchronized methods (or nested synchronized blocks).

Practice

- Review an incorrect bank account examples that synchronizes the thread's method, not the Bank Account methods.
- Review the deadlock example
 - MyObject, DeadlockThread

VARIABLES

Atomic

- An *atomic* action happens all at once
 - It cannot stop in the middle
 - It either happens completely or not at all
- `i++` is **not** atomic- there are multiple steps

Atomic Variables

- `java.util.concurrent.atomic` contains classes that support atomic actions on variables
- Example: `AtomicInteger`
 - `get`
 - `set`
 - `incrementAndGet`
 - `decrementAndGet`
 - `updateAndGet(IntUnaryOperator)` (`IntUnaryOperator`: param int, return int)
- Using an atomic variable can help avoid unnecessary synchronization

Practice

- Review the revised Counter example that uses atomic integer
 - AtomicCounter, AtomicTest

LongAdder

- Java 8 added a new class LongAdder designed represent a collective sum.
- Methods
 - add(long)
 - increment
 - sum

volatile

- If one thread changes the value of a variable, you expect other threads to see that change.
- But that's not always the case!
 - The JVM allows values to be stored temporarily in working memory (like a cache), rather than main memory (the heap).
 - There is no guarantee that the heap version has been updated- the updated value could still be stored in cache.
- So one thread could change the value and that change is stored in cache. But then another thread reads from the heap and gets the wrong value.

volatile

- Declaring a variable volatile enforces that all updates will be written back to heap memory immediately.
- It also ensure that reads/writes are atomic (all at once).
 - Reads and writes are atomic for reference variables and for most primitives (NOT for long and double).
 - Reads and writes are atomic for *all* variables declared volatile (including long and double).

volatile- NOT synchronized!!

- volatile does **not** mean atomic overall!
 - myVolatileVariable++ is still not atomic!
 - volatile is **not** the same as synchronized
- volatile may be good enough if you have only one thread updating the variable and many others only reading it
 - if you have multiple threads updating, you need to synchronize!

WAIT/NOTIFY

Locks

- Every object has an *intrinsic lock*.
 - Also called a *monitor lock* or just *monitor*.
- A thread can *acquire* a lock.
 - It then *owns* the lock.
 - As long as it owns the lock, no other thread can acquire the lock.
 - When the thread is done, it can *release* the lock.

Synchronization

- If thread A tries to enter a synchronized block and thread B is already executing that block (meaning it *holds the lock on the synchronized object*), thread A is placed into a waiting area.
- When thread B finishes, it releases the lock.
- The JVM then removes A from the waiting area and assigns the lock to thread A.
- This is one specific application of the wait/notify mechanism.

wait/notify/notifyAll

- A thread *waits* on a condition to continue executing.
- Some other thread updates that condition (*notifies the thread*) when it can execute.
- notify wakes up a single thread (decided by the ThreadScheduler)
 - We don't automatically know which one will be notified- only good for "massively parallel" operations where all threads doing the same job.
- notifyAll wakes up *all* threads

Waiting and Notifying

```
synchronized(object) {  
    while(conditionThatMustBeTrue) {  
        try {  
            object.wait();  
        } catch(InterruptedException ex) {}  
    }  
    object.doTheWork();  
}
```

```
synchronized(object) {  
    // work is ready to be done  
    conditionThatMustBeTrue = true;  
    object.notifyAll();  
}
```

Waiting

- You have to own the lock in order to invoke wait.
 - Putting the wait() call inside a synchronized block is one way to ensure this.
 - Invoking wait() releases the lock
 - When the thread is notified, it will have to continue waiting until the lock is released
- Always wait inside of a loop that checks the condition.
 - There is no guarantee that the condition has changed as a result of the notification.
 - So you want to recheck it before moving on.

Practice

- Review the deposit/bonus methods of the bank account classes.
 - Multiple threads will make random deposits.
 - There are “bonus” threads that wait to add a \$1 “savings bonus” when the balance passes a certain amount.
 - Classes: WNBankAccount, WNBankThread, WNBankTester

Producer/Consumer Relationship

- An application that shares data between two threads: a producer and a consumer
- A producer thread creates some data/item/object that a consumer thread uses.
 - Example: read an object from a database and pass it off for processing
 - Example: read input from a user and hand it off to be analyzed
- Coordination is essential!
- The consumer cannot use item unless it's been produced.
 - It has to wait.
- The producer cannot create an item if the consumer is still busy or hasn't retrieved the old item.
 - It has to wait.

Practice

- Review the producer/consumer example.
 - The *number box* holds a single number and can either be filled or empty.
 - Classes: `ProducerThread`, `ConsumerThread`, `ProducerConsumerTest`, `NumberBox`

Using wait/notify

- Most people don't directly use wait/notify anymore.
- As Joshua Block (Effective Java) states:
 - “using wait and notify directly is like programming in “concurrency assembly language,” as compared to the higher-level language provided by `java.util.concurrent`. **There is seldom, if ever, a reason to use wait and notify in new code.**”
- Instead, use high-level constructors, including `BlockingQueue`.

Blocking Queue

- A special queue that supports the producer/consumer pattern.
 - Threads must wait for a queue to be non-empty to retrieve.
 - Threads must wait for space to add.
- Different methods based on the action you want to take when an add/remove cannot be handled immediately.

Summary of BlockingQueue methods

	<i>Throws exception</i>	<i>Special value</i>	<i>Blocks</i>	<i>Times out</i>
Insert	<code>add(e)</code>	<code>offer(e)</code>	<code>put(e)</code>	<code>offer(e, time, unit)</code>
Remove	<code>remove()</code>	<code>poll()</code>	<code>take()</code>	<code>poll(time, unit)</code>
Examine	<code>element()</code>	<code>peek()</code>	<i>not applicable</i>	<i>not applicable</i>

Practice

- Review the producer/consumer example using a BlockingQueue.
 - `ProducerThreadBQ`, `ConsumerThreadBQ`, `ProducerConsumerTestBQ`

LOCKS

Locks

- A *lock* ensures that only one thread can access a block of code at a time.
- Use the Lock interface and ReentrantLock class

```
myLock.lock();  
try {  
    // code to be protected  
} finally {  
    myLock.unlock(); // must be in finally!  
}
```

Locks

- Locks can provide similar functionality to synchronized, and they have internal wait/notify mechanisms built in.
- They also have additional benefits:
 - tryLock method (backs out if a lock is not available)
 - lockInterruptibly (backs out if interrupted while waiting)
 - Better performance under high *thread contention*
 - Can include multiple condition variables
 - Can split locking across methods (e.g., get a lock in one method and release it in another)

Practice

- Review the example of the BankAccount with fee and bonus, using locks instead of synchronized
 - Classes: LBankAccount, LBankThread, LBankTester

Conditions

- You can add multiple conditions to a lock.
 - This can replicate wait/notify, but is more flexible since you can have multiple conditions.

```
private Condition myCondition;
```

```
myCondition = myLock.newCondition();
```

```
myCondition.await();
```

```
myCondition.signal(); // or signalAll()
```

Locks and Conditions

```
private Lock myLock = new ReentrantLock();
private Condition myCondition = myLock.newCondition();

myLock.lock();
try {
    while (!someConditionThatMustBeTrue) {
        myCondition.await ();
    }
    // do the work that requires the condition and lock
} catch (InterruptedException e) {}
} finally {
    myLock.unlock();
}
```

```
myLock.lock();
try {
    change someConditionThatMustBeTrue;
    myCondition.signalAll();
} finally {
    myLock.unlock();
}
```

Practice

- Review the BankAccount deposit/bonus example that uses locks and conditions.
 - A bonus is given only when the balance is over \$5000.
 - Classes: LCBankAccount, LCBankThread, LCBankTester

Practice

- Write a program to simulate the one-lane bridge problem using locks and conditions.
- There is a one-lane bridge where a car can only cross in one direction at a time.
- Write a simple bridge that doesn't consider the cars' directions.
 - Any car can go next, just one car at a time.
- Write a bridge that does take direction into .
 - The direction of the current car is given priority for who gets to go next.

EXECUTORS AND THREAD POOLS

Executors

- A downside of low-level thread code is that it links our task to our task-execution policy (e.g., create one thread for every task).
- Also, it is **expensive** to create new threads.
- *Executors* allow you to run multiple threads without manually managing them.
 - They also re-use threads, rather than always creating new ones.

Thread Pools

- Thread pools are groups of threads, known as *worker threads*.
 - Worker threads sit idle and ready to run.
- Using a thread pool can reduce the expense of creating threads.
 - Important for large-scale applications
- Tasks are submitted to the pool and queued up if all threads in the pool are busy.
- Example:
 - Consider a website that can handle 10 requests at a time. (Hope they don't get too popular!)
 - Without a thread pool, when the 11th request comes in, an 11th thread is created, and the system then freezes- for **all** threads.
 - With a pool, the 11th task is queued up and waits for one of the 10 threads to be free.

Using Executors

- ExecutorService object created through factory method of the Executors class:

```
ExecutorService executor =
```

```
    Executors.newSingleThreadExecutor();  
    // all tasks execute on a single thread
```

or

```
    Executors.newFixedThreadPool(poolSize);  
    // all tasks execute on pool of threads
```

or

```
    Executors.newCachedThreadPool();  
    // all tasks execute on pool of threads
```

Using Executors

- Execute the executor

```
executor.execute(Runnable r);
```

```
executor.execute( () -> runnable code );
```

- You can call execute multiple times for repeated tasks.
- You can also send in an object of type `Callable<V>`
 - Callable is analogous to Runnable, but it returns a value!
- Stop the executor (must do!!):

```
executor.shutdown();
```

Practice

- Review the dice example to compare low-level thread control to executor control.
 - Class: DiceTester

IMMUTABLE OBJECTS

Immutable Objects

- An object is *immutable* if its state cannot change once it is constructed.
- Example- the String class
 - `String s = "Jessica";`
 - `s.toUpperCase();` // does not change s!
 - `String s2 = s.toUpperCase();`
 - `s = s.toUpperCase();`

Immutable Objects

- An object is *immutable* if its state cannot change once it is constructed.
- Immutable objects are great for use in multi-threaded programs.
- Because their state cannot change, you don't have to be worried about corruption through thread interference or inconsistent states.
- Immutable objects are thread-safe and can be used outside of synchronized blocks.

Objects in Memory- Modifying an Object


- NOT POSSIBLE with an immutable object!

```
MyObject ref1 = new MyObject(...);
```

```
MyObject ref2 = ref1;
```

```
ref1.setVariable2(...);
```

```
// not allowed if MyObject is immutable!
```




Objects in Memory- Changing a *Reference*

- Still allowed with an immutable object! You are not changing an object!

```
MyObject ref1 = new MyObject(...);
```

```
MyObject ref2 = ref1;
```

```
ref1 = new MyObject(...);
```



Objects in Memory

- Invoking a method with the **dot operator** most likely **changes the object itself**
 - The state or characteristic of the actual object changes.
- Assigning a new value **wchanges what the variable/reference points to**
- ith the **equals sign** just This “changes the memory arrows” but does not change the actual object

How to Create Immutable Classes

- Make the class final
 - Or make the constructor private and provide static factory method
- No setter methods
- Make all instance data “variables” final and private
- For instance variables that are **mutable** objects:
 - Ensure there are no methods that would modify those objects
 - Do not share references to those objects!
 - When passed an object, make a copy before storing it as instance data
 - Example: in a constructor!
 - When returning an object, make a copy to return
 - Example: in getter methods!

Practice

- Review the EmployeeIM example.

COLLECTIONS

Multiple Threads and Collections

- It is easy to corrupt a data structure if accessed by multiple threads.
- You should not use a collection with no thread support in a concurrent program.
- Instead, use a:
 - Synchronized Collection
 - Concurrent Collection

Practice

- Look at an example of filling and emptying a list with threads using a normal non-thread-safe collection.
 - Classes: ListAdder, ListRemover, ConcurrentCollectionsTest

Synchronized Collection Classes

- All **methods** in the class are synchronized.
- Classes
 - Vector
 - Hashtable (**not** HashMap or TreeMap)
- Create your own (preferred to above):
 - `Collections.synchronizedList(myList)`
 - `Collections.synchronizedMap(myMap)`
 - `Collections.synchronizedSet(mySet)`

Practice

- Re-run the add/remove list example with a synchronized list.

Thread Safety- Important!!

- Using a synchronized collection assures no corruption from invoking a method.
- But it doesn't mean you're off the hook for thread safety overall!
- **Combinations of methods are not atomic!**

- Example:

```
if (!testList.isEmpty())  
    testList.remove(0);
```

- Fix:

```
synchronized(testList) {  
    if (!testList.isEmpty())  
        testList.remove(0);  
}
```

Practice

- Re-run the add/remove list example with a synchronized list AND thread-safe code.

Iterators

- A `ConcurrentModificationException` is thrown if one thread modifies a synchronized collection while an iterator is retrieving an item.
- You'd have to synchronize access to the block of code iterating to prevent this.
 - Costly!

Concurrent Collection Classes

- More sophisticated support
 - Thread safe and scalable!
- These classes allow concurrent access to different parts of a single data structure.
- Classes
 - ConcurrentHashMap
 - More efficient than `Collections.synchronizedMap(myMap)`
 - ConcurrentLinkedQueue

Iterators

- The concurrent collection classes return *weakly consistent iterators*
- These iterators will **not** throw a `ConcurrentModificationException`
- The iterator may or may not reflect all modifications made since construction
 - It will **not** return an item twice
 - If an element was removed before iteration starts, it won't be returned
 - If an element is added after iteration starts, it may or may not be returned

ConcurrentHashMap

- Retrieval operations do not block
- Supports multiple concurrent write operations
- Use traditional Map methods, plus **atomic versions** of compound operations:
 - putIfAbsent(key, value)
 - compute(key, BiFunction<K, V> remapFunction)
 - BiFunction: parameter K and V, returns V
 - computeIfPresent(key, BiFunction)
 - computeIfAbsent(key, BiFunction)

ConcurrentHashMap

- Bulk operations
 - search- search elements
 - reduce- accumulate elements
 - forEach- perform an action on each element
- Each can be applied to Keys, Values, both, or Map.Entry
- Each takes a parallelismThreshold and a function or bi function
 - Threshold: if the map has more elements than the threshold, it will be parallelized

Thread Safety- Important!!

- Again, using a concurrent collection only assures no corruption from invoking a method.
- But it doesn't mean you're off the hook for thread safety overall!
- Combinations of methods are still not atomic!

Practice

- Review the scrabble words example.
 - Make a map of the number of words that start with each letter.
 - Make a map of the sum of the score of words that start with each letter.
 - Classe: LetterStartThread, LetterStartThreadSafe, WordScoreSumThreadSafe, ConcurrentMapTestingScrabble